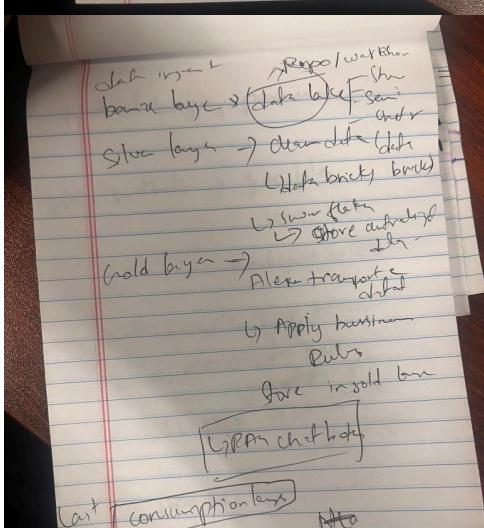
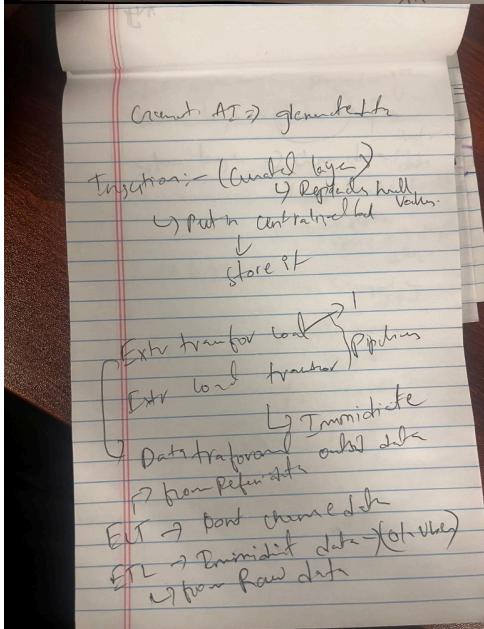
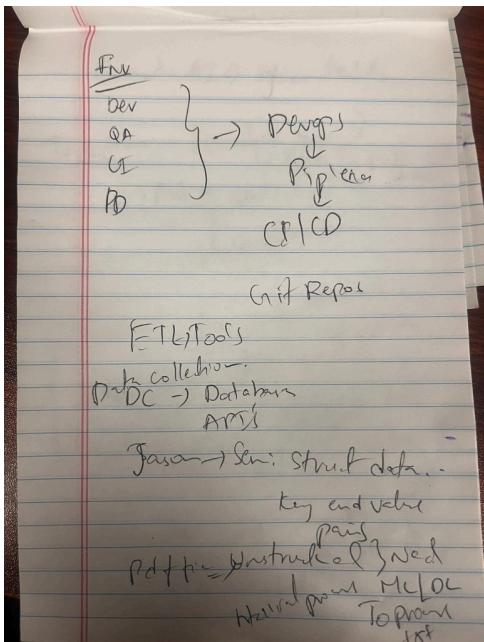


# Day-1

Wednesday, January 21, 2026 7:24 PM



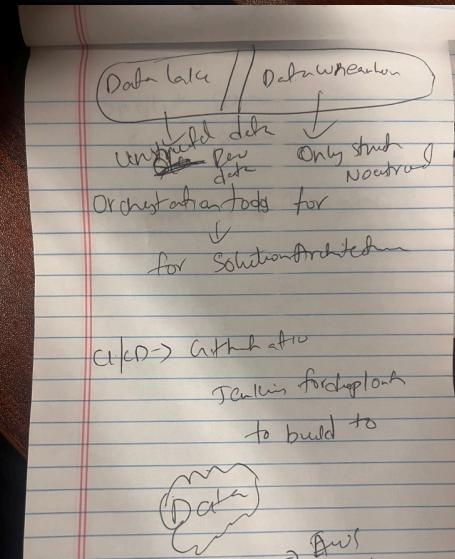
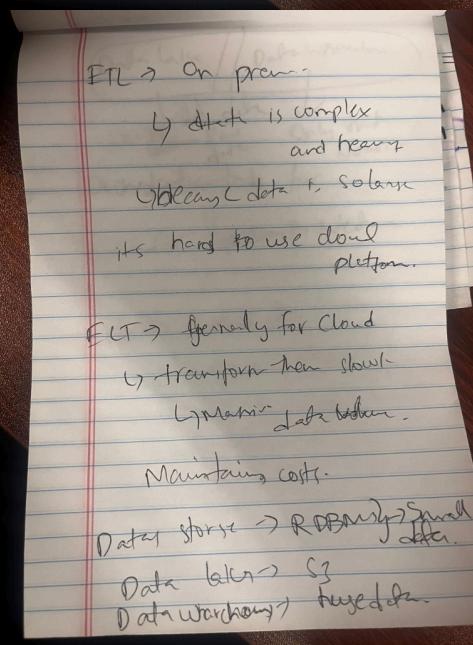
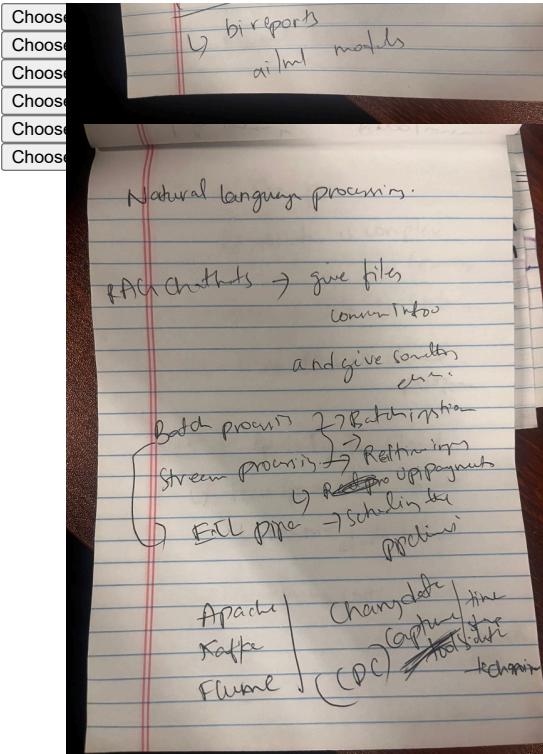
- “Every data has a story to tell.”
- The role of:
  - Data Engineers** → build pipelines & infrastructure
  - Data Scientists / Analysts** → extract insights and meaning
- Your job is to give voice to data through processing and analysis.

- Data Collection**
  - Databases: MySQL, PostgreSQL, MongoDB
  - APIs
  - Files: CSV, JSON, XML
  - Web scraping (BeautifulSoup, Scrapy)
  - Event streams
- Data Types**
  - Structured** → relational tables
  - Semi-structured** → JSON, XML
  - Unstructured** → PDFs, images, videos
- Data Ingestion**
  - Moving data into centralized storage
  - Methods:
    - ETL** (Extract → Transform → Load)
    - ELT** (Extract → Load → Transform)

- Batch Processing**
  - Scheduled runs (daily, weekly, monthly)
  - Used for reports, analytics, historical data
- Stream / Real-Time Processing**
  - Instant processing (Uber, UPI, food delivery)
  - Tools mentioned: **Kafka**, **Apache NiFi**
- CDC (Change Data Capture)**
  - Captures only changed data using timestamps
  - Efficient for real-time updates

- Bronze Layer (Data Lake)**
  - Raw, unprocessed data
- Silver Layer**
  - Cleaned and processed data
- Gold Layer**
  - Business-ready data with rules applied

- Databases**
  - Smaller, structured data
- Data Warehouses**
  - Large historical data
  - Uses **fact & dimension tables**
  - Star schema / Snowflake schema**



- **Data Lakes**
- Raw, semi-structured & unstructured data

- **Power BI dashboards**
- Reporting & analytics
- Machine Learning models
- Predictive analytics
- NLP & Generative AI use cases

- **Data Governance**
- Security, privacy, compliance
- Data masking, NDAs, sensitive data handling
- **DevOps**
- CI/CD pipelines
- Managing dev → test → prod environments
- **Orchestration**
- Tools like **Airflow, Azure Data Factory (ADF)**
- Schedule & monitor pipelines

Data warehouses are designed to use parallel processing so that large volumes of data can be processed faster by dividing the workload across multiple CPUs or nodes simultaneously.

#### Parallel processing

Instead of:

- 1 CPU → slow
- You get:
- Many CPUs / nodes → fast

ETL pitel / spark

Shuttle Schema

Strav Schema

Anyflow Prefed/Dagster

J  
Orchestration

DB (AWS Cloud)

ETL

Mater agit Repo  
should call the  
data.